

Mean Squared Error (MSE)

Formula

$$J_{\text{MSE}} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

where:

- y_i are the true target values,
- \hat{y}_i are the predicted values,
- n is the number of observations,
- J_{MSE} denotes the MSE-based loss function (here we use J to represent any loss/cost function).

Best for

- **Penalizing large errors heavily:** Squaring the differences means that larger errors have a disproportionately higher impact on the overall cost.

Characteristics

- **Differentiable and easy to compute:** The MSE-based cost function is smooth and lends itself well to gradient-based optimization.
- **Sensitive to outliers:** Squaring amplifies large errors, making the metric more sensitive to outliers.
- **Encourages predictions close to the mean:** Especially when the target distribution is unimodal, the MSE criterion drives predictions toward the mean of the targets.

Interpretations

Gaussian Distribution

The probability density function (PDF) of a Gaussian (Normal) distribution with mean μ and variance σ^2 for a variable x is:

$$\mathcal{L}(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right),$$

where we use \mathcal{L} to indicate the likelihood function.

Probabilistic Perspective (Derivation)

Assuming the prediction errors follow a Gaussian distribution with mean 0 and variance σ^2 , the likelihood (\mathcal{L}) for a single observation y_i given the predicted value \hat{y}_i is:

$$\mathcal{L}(y_i | \hat{y}_i) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - \hat{y}_i)^2}{2\sigma^2}\right).$$

1. *Write out the PDF:*

$$\mathcal{L}(y_i | \hat{y}_i) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - \hat{y}_i)^2}{2\sigma^2}\right).$$

2. *Take the logarithm:*

$$\log \mathcal{L}(y_i | \hat{y}_i) = \log\left(\frac{1}{\sqrt{2\pi\sigma^2}}\right) + \log\left(\exp\left(-\frac{(y_i - \hat{y}_i)^2}{2\sigma^2}\right)\right).$$

3. Use log rules:

$$\log \mathcal{L}(y_i | \hat{y}_i) = -\frac{1}{2} \log(2\pi\sigma^2) - \frac{(y_i - \hat{y}_i)^2}{2\sigma^2}.$$

4. Multiply by -1 to define the negative log-likelihood (our cost J):

$$J_i = -\log \mathcal{L}(y_i | \hat{y}_i) = \frac{(y_i - \hat{y}_i)^2}{2\sigma^2} + \frac{1}{2} \log(2\pi\sigma^2).$$

Single-Point vs. Full Likelihood It's common to discuss a single data point's probability density as a likelihood in a derivation. Formally, the word likelihood typically refers to the joint function of all data points, but each individual term in that product (or each individual summand in the log-likelihood) can also be referred to as the likelihood contribution of a single observation. Thus, we can show the negative log-likelihood derivation on a per-data-point basis, then multiply (or sum in log space) to account for the entire dataset.

Negative Log-Likelihood for the Entire Dataset

For n independent observations, the joint likelihood is the product of individual likelihoods. Consequently, the negative log-likelihood for all observations becomes the sum:

$$J_{\text{NLL}} = -\log \mathcal{L}(\{y_i\} | \{\hat{y}_i\}) = \sum_{i=1}^n \left[\frac{(y_i - \hat{y}_i)^2}{2\sigma^2} + \frac{1}{2} \log(2\pi\sigma^2) \right].$$

Note on Curly Braces ($\{\}$): When we write $\{y_i\}$ or $\{\hat{y}_i\}$, we're indicating the entire set (or collection) of y_i or \hat{y}_i values for $i = 1, \dots, n$. This is standard mathematical notation for describing a group of elements, rather than a single element.

Ignoring the constant term $\frac{1}{2}n \log(2\pi\sigma^2)$ gives:

$$J_{\text{NLL}} = \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + (\text{const.}).$$

A Note on the $\frac{1}{2\sigma^2}$ Factor

From a purely optimization perspective (where σ^2 is fixed and we only optimize \hat{y} or model parameters), this constant factor does not affect the location of the minimum. So, many treatments omit it when they only care about the solution for \hat{y} . However, when the variance σ^2 is also being optimized (as in a full maximum-likelihood approach for both mean and variance), we need to keep that factor to reflect the exact Gaussian log-likelihood.

Minimizing this with respect to \hat{y}_i is equivalent to minimizing the sum of squared errors. Hence, **minimizing the negative log-likelihood** under these Gaussian assumptions is equivalent to **minimizing MSE**.

Geometric Perspective (Derivation)

The MSE can also be interpreted as the squared Euclidean (L2) distance between the predicted and true values:

$$\|y - \hat{y}\|_2^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2.$$

Minimizing this distance is equivalent to finding the point \hat{y} closest (in an L2 sense) to the actual target y .

Connection to Linear Regression

MSE is the standard cost function in Ordinary Least Squares (OLS) regression. In linear regression, the model is often expressed as:

$$y = X\beta + \varepsilon,$$

where X is the design matrix of input features, β is the parameter vector, and ε represents normally distributed noise. Minimizing the MSE cost:

$$J_{\text{MSE}}(\beta) = \sum_{i=1}^n (y_i - X_i\beta)^2$$

yields the OLS estimates. If the columns of X are linearly independent, the closed-form solution is:

$$\hat{\beta} = (X^T X)^{-1} X^T y.$$

Detailed Derivation of the OLS Solution

Here is a step-by-step outline of how we arrive at the formula $\hat{\beta} = (X^T X)^{-1} X^T y$ in linear algebra terms:

1. Set up the cost function:

$$J_{\text{MSE}}(\beta) = \sum_{i=1}^n (y_i - X_i\beta)^2.$$

In matrix form, if y is the $n \times 1$ vector of targets and X is the $n \times d$ matrix of features (with each row corresponding to one observation), then:

$$J_{\text{MSE}}(\beta) = (y - X\beta)^T (y - X\beta).$$

2. Take the gradient w.r.t. β : We can expand the above cost:

$$(y - X\beta)^T (y - X\beta) = y^T y - y^T X\beta - (X\beta)^T y + (X\beta)^T (X\beta).$$

Notice that $y^T X\beta$ and $(X\beta)^T y$ are scalars (single numbers), so they are equal. Thus:

$$(y - X\beta)^T (y - X\beta) = y^T y - 2y^T X\beta + \beta^T X^T X\beta.$$

Now take the derivative (gradient) w.r.t. β :

$$\nabla_{\beta} J_{\text{MSE}}(\beta) = -2X^T y + 2X^T X\beta.$$

(We use basic rules of matrix calculus here, noting that the derivative of $\beta^T A\beta$ w.r.t. β is $2A\beta$ if A is symmetric. Here, $X^T X$ is symmetric.)

3. Set the gradient to zero (the Normal Equations): To minimize $J_{\text{MSE}}(\beta)$, we set its gradient to 0:

$$-2X^T y + 2X^T X\hat{\beta} = 0 \quad \Rightarrow \quad X^T X\hat{\beta} = X^T y.$$

This equation is known as the *Normal Equation*.

4. Solve for $\hat{\beta}$: Provided $X^T X$ is invertible (which requires that X has full column rank), we can multiply both sides by $(X^T X)^{-1}$:

$$\hat{\beta} = (X^T X)^{-1} X^T y.$$

5. Interpretation:

- $X^T X$ captures the correlations among features.
- $X^T y$ captures how each feature relates to the target vector.
- $(X^T X)^{-1} X^T y$ thus gives the coefficients that minimize the overall sum of squared errors.

This derivation relies on understanding how to take matrix derivatives and the condition that $X^T X$ be invertible. When these assumptions hold, the formula neatly expresses the solution that best fits the data in the least squares sense.

Summation-Based Derivation (Without Matrix Algebra)

While matrix notation is concise, we can also derive the same result using summations explicitly. Let's assume our model is:

$$\hat{y}_i = \beta_0 + \sum_{j=1}^d \beta_j x_{ij},$$

where x_{ij} is the value of the j -th feature for the i -th data point, β_0 is an intercept term, and β_j are the feature coefficients.

Our goal is to minimize the MSE cost function:

$$J_{\text{MSE}}(\beta_0, \beta_1, \dots, \beta_d) = \sum_{i=1}^n (y_i - \hat{y}_i)^2.$$

Substituting \hat{y}_i :

$$J_{\text{MSE}} = \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^d \beta_j x_{ij} \right)^2.$$

1. Take partial derivatives w.r.t. each β_k For β_0 :

$$\frac{\partial J_{\text{MSE}}}{\partial \beta_0} = \sum_{i=1}^n -2 \left(y_i - \beta_0 - \sum_{j=1}^d \beta_j x_{ij} \right) \cdot 1 = 0$$

(when set to 0 for the optimum). Simplify to get:

$$\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^d \beta_j x_{ij} \right) = 0.$$

For each β_k with $k \geq 1$:

$$\frac{\partial J_{\text{MSE}}}{\partial \beta_k} = \sum_{i=1}^n -2 \left(y_i - \beta_0 - \sum_{j=1}^d \beta_j x_{ij} \right) x_{ik} = 0.$$

2. System of $d+1$ equations These partial derivatives give us $d+1$ simultaneous equations (one for β_0 , and one for each β_k , $k = 1, \dots, d$). In compact form:

$$\begin{cases} \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^d \beta_j x_{ij}) = 0, \\ \sum_{i=1}^n x_{i1} (y_i - \beta_0 - \sum_{j=1}^d \beta_j x_{ij}) = 0, \\ \vdots \\ \sum_{i=1}^n x_{id} (y_i - \beta_0 - \sum_{j=1}^d \beta_j x_{ij}) = 0. \end{cases}$$

Solving this system yields the same result as applying the matrix-based Normal Equation: each equation here corresponds to a row in $X^T X \beta = X^T y$.

Hence, whether we choose the summation approach (manually solving these $d+1$ equations) or the matrix approach (using linear algebra), we arrive at the same solution. The matrix form is simply a more compact and elegant representation of these summations.

Gradient-Based Optimization

Because J_{MSE} is smooth and differentiable, it fits perfectly in gradient-based methods. The partial derivative of J_{MSE} with respect to the prediction \hat{y}_i is:

$$\frac{\partial J_{\text{MSE}}}{\partial \hat{y}_i} = \frac{2}{n}(\hat{y}_i - y_i).$$

When using a parametric model, this gradient is propagated back through the model parameters (e.g., weights in neural networks) by the chain rule.

Relationship to R^2

Minimizing J_{MSE} is closely related to maximizing the coefficient of determination R^2 , which measures how well the model explains the variance in the observed data. The R^2 value is defined (in one of its forms) as:

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2},$$

where \bar{y} is the mean of the true targets y_i . Reducing $\sum_{i=1}^n (y_i - \hat{y}_i)^2$ (i.e., the SSE) will increase R^2 . A higher R^2 indicates a tighter fit of the model to the data.